



COSMO CONSULT

Business-Software for People

DYNAMICALLY SCALING BC CONTAINERS ON DOCKER SWARM USING
AZURE SQL AND TRAEFIK

AREOPA WEBINAR BY TOBIAS FENSTER, DEC 10, 2019



Tobias Fenster

CTO at COSMO CONSULT Group

Dual Microsoft MVP for Business
Applications and Azure

🐦 @tobiasfenster

📡 tobiasfenster.io

✉ tobias.fenster@cosmoconsult.com

in tobiasfenster



COSMO CONSULT

Business-Software for People



WHAT ARE WE TALKING ABOUT:

SCOPE AND CHALLENGE

Starting point



- You are running or want to run your dev or test NAV / BC environments in Docker containers
- You want to run multiple containers on the same VM
- You have solved the networking challenges (if not, see the Areopa webinar of Oct 1st 2019)

Challenges / limitations


- You need more resources than one VM can provide
- You want to dynamically scale up and down
- You don't want to run SQL server in the same containers as NAV / BC (not very efficient, SQL Express, same scaling limits)



WHAT ARE WE TALKING ABOUT: SOLUTION



PART 1 OF THE SOLUTION: DOCKER SWARM




- Built-in **container orchestrator** from Docker
- Main benefits / features:
 - Brings resources of **multiple container hosts** together
 - **Central** management and control
 - Share **configuration and secrets** across the swarm
- **Declarative** service model
- Automatic **"self-healing"** concepts
- Advanced networking for **resiliency**
- What about **Kubernetes** and the recent **Mirantis** deal?



PART 1 OF THE SOLUTION:

DOCKER SWARM



- Some **basics** in the Docker Swarm world:
 - **Service** = declaration of the images, number of containers (tasks) and configurations you want to run; can be replicated or global
 - **Node** = container host / engine that joined the swarm
 - **Manager** = node with control rights
 - **Worker** = node that executes tasks


- How to “run” something:
 1. **Declare** your service
 2. **Submit** that to a manager node
 3. **Swarm** creates necessary **tasks on nodes** and keeps **desired state**



PART 1 OF THE SOLUTION:

DOCKER SWARM


Goal: 5 containers
on 3 hosts



PART 1 OF THE SOLUTION:

DOCKER SWARM


Goal: 5 containers
on a **Swarm** with
3 hosts



PART 1 OF THE SOLUTION:

DOCKER SWARM


Goal: Connect to the
service, no matter
which task



PART 1 OF THE SOLUTION:

DOCKER SWARM

Goal: Connect to the
service, no matter
which task




PART 1 OF THE SOLUTION:

DOCKER SWARM

Goal: Connect to the
service, no matter
which task




Doesn't directly work with BC
as it requires **stateless service**



PART 1 OF THE SOLUTION:

DOCKER SWARM


Goal: Services are able to find other services, no matter which host




PART 1 OF THE SOLUTION:

DOCKER SWARM

Goal: Services are able to find other services, no matter which host



PART 1 OF THE SOLUTION: DOCKER SWARM



- Demo scenario:

- Create 1 service consisting of 2 containers

- Check connections

- Scale service up

- See placement on nodes



- Remove task

- See recovery

→ Let's see it!




WHAT ARE WE TALKING ABOUT: SOLUTION



PART 2 OF THE SOLUTION:

AZURE SQL WITH AN ELASTIC POOL



- Platform as a Service (PaaS) offering
 - Always **kept current** by Microsoft
 - **Scale** up and down **dynamically**
 - (Almost) **no resource limits**
 - Elastic pools allow **resource sharing** across all databases (with configurable limits)



→ No server maintenance, no limits; just cost...



WHAT ARE WE TALKING ABOUT: REMEMBER THE START?



Challenges / limitations

- You need **more resources** than one VM can provide
- You want to **dynamically scale** up and down
- You don't want to run **SQL server in the same containers** as NAV / BC (not very efficient, SQL Express, same scaling limits)



BRINGING IT TOGETHER: SWARM (AND TRAEFIK) & AZURE SQL

- Prereq: Azure SQL database with a “template” database
 - 1 ARM template to deploy
 - 1 **manager** (SPOF!) and x **workers**
 - 1 Azure SQL server with 1 **elastic pool**
 - Setup script to **initialize** the Swarm on the manager, **join** it on the workers and set up **Traefik**
 - Script to prepare **access credentials** for template DB and target pool and share them as **secrets**
 - Specific scenario: Start a **BC Swarm service** and connect to a **DB created on demand** as copy of the template
- Let's see it





COSMO CONSULT

Business-Software for People

Get **production support** for Docker containers
with D365 BC: <http://bit.ly/DockerProd>

Get rid of the **120kB size limit** for dev
licenses: <http://bit.ly/UnlimitedFlf>

THANK YOU

FOR YOUR TIME

WHICH QUESTIONS CAN I
TRY TO ANSWER?

